

VISHMITHA POOJARY

Mangalore, Karnataka ☎ +91 9110833308 ✉ vishuvishmitha84@gmail.com in LinkedIn 🌐 GitHub

PROFESSIONAL SUMMARY

Software engineer specializing in backend systems, APIs, and AI-powered applications. Uses AI coding tools (Codex, Claude, Gemini, GPT) as a structured part of the development workflow — creating SKILL.md reference files and using /plan to generate implementation plans before coding, reviewing AI output critically, and when AI produces repeated mistakes, diagnosing the root cause and providing the right context to correct it rather than re-prompting blindly. Treats AI as a reliable development partner that speeds up building while keeping engineering decisions sound.

EDUCATION

Course	Institution	Year	CGPA / %
MCA	St Aloysius Institute of Management	2024–2026	9.00
BCA	SDM College of Business Management	2021–2024	9.15

SKILLS

AI Dev Workflow: Codex, Claude, Gemini, ChatGPT — SKILL.md, /plan, structured prompting, AI-assisted debugging and root-cause analysis

AI & Agents: LangGraph, Gemini API, OpenAI API, Prompt Engineering, Agentic Function Calling, Multi-Agent Orchestration

Backend: Python, FastAPI, Node.js, TypeScript, Express.js, Hono.js **Frontend:** React.js, Next.js, TailwindCSS

Databases & Tools: PostgreSQL, MongoDB, Drizzle ORM, SQLAlchemy, Zod, JWT, Docker, Git, WebSockets, Pydantic

EXPERIENCE

Software Engineer — Trikon Software Labs <i>Mangalore, India</i>	Jul 2026 – Present
– Build and maintain backend services for a live enterprise voice AI platform (Node.js, TypeScript, Express, PostgreSQL, Gemini Live API) integrating telephony providers (Plivo, Vobiz) via real-time WebSocket audio streaming – Replaced static pre-recorded greetings with a dynamic multilingual AI voice engine supporting 30+ languages; built audio filters to stop AI getting confused by background noise, preventing hallucinations and runaway calls – Rebuilt the outbound dialer from a polling-based system to an event-driven architecture with smart traffic pacing — the system now handles high-volume campaigns without exceeding telecom rate limits or crashing under load – Executed a zero-data-loss live production database migration (Supabase → Neon PostgreSQL) using a two-step strategy — schema recreation then raw data extraction — that bypassed failures standard migration tools could not handle – Use Codex and Claude daily with structured SKILL.md context files and /plan to generate implementation plans before each feature; review plans before coding and when AI produces repeated errors, trace the actual execution flow to identify the root cause and re-prompt with the right context rather than guessing	
Software Developer Intern — TechBrew Solutions <i>Mangalore, India</i>	Jan 2026 – Jun 2026
– Hospital Laundry Management System: Built production REST APIs end-to-end from OpenAPI contracts (Hono.js, Drizzle ORM + dbmate, Zod, PostgreSQL) for a live hospital client — schema design, migrations, validation, and React frontend integration; collaborated via GitHub PRs in a Docker environment – AI Multi-Agent Trading System: Worked on a LangGraph-based 4-agent system (Supervisor, StockNews, TechnicalAnalysis, TradeBrain) — used /plan and Claude to design the agent architecture and message-passing flow before implementation; engineered structured LLM prompts with event taxonomy, confidence/impact scoring, and Pydantic-validated JSON output – Used AI tools throughout both projects: fed actual errors and execution traces to Claude/GPT when debugging rather than re-prompting blindly, and used SKILL.md files to give AI the right codebase context upfront	

PROJECTS

AI Hospital Voice Receptionist (Deployed) FastAPI, PostgreSQL, Gemini, LiveKit, Deepgram, ElevenLabs	2026
– Designed and deployed an agentic AI voice system for healthcare automating appointment booking and rescheduling in English and Hindi; used Codex with a SKILL.md to scaffold the agentic function-calling architecture and debug LiveKit WebRTC integration issues systematically – Integrated LiveKit (WebRTC), Deepgram (STT), and ElevenLabs (TTS) with Gemini function calling against live PostgreSQL availability data for low-latency, human-like voice interactions	
AI Interview Platform Backend (Ongoing) FastAPI, PostgreSQL, SQLAlchemy, JWT, LLM APIs	2026
– 9 secured REST endpoints — JWT auth, LLM-driven adaptive question generation, PDF resume skill extraction, weak-skill tracking, session persistence, and leaderboard analytics; planned the full API structure using /plan before writing any code	
AI Multi-Agent Intraday Trading System Python, LangGraph, Gemini API, OpenAI API, Pydantic	2026
– Engineered structured LLM news ranker with event taxonomy (earnings, M&A, regulatory, debt restructuring), confidence/impact/freshness scoring, and Pydantic-validated JSON output — iteratively refined by comparing Gemini and GPT outputs to find which model handled edge cases better – Fixed a cross-batch deduplication bug where the same article appeared under multiple tickers, bypassing the global article limit — traced root cause through the pipeline, solved by merging all batches before deduplication and applying global TOP-N post-merge; built 25s timeout, 3-retry exponential backoff, and per-call token/latency tracking	